Lecture - 8

⊗ Bias | Variance :-

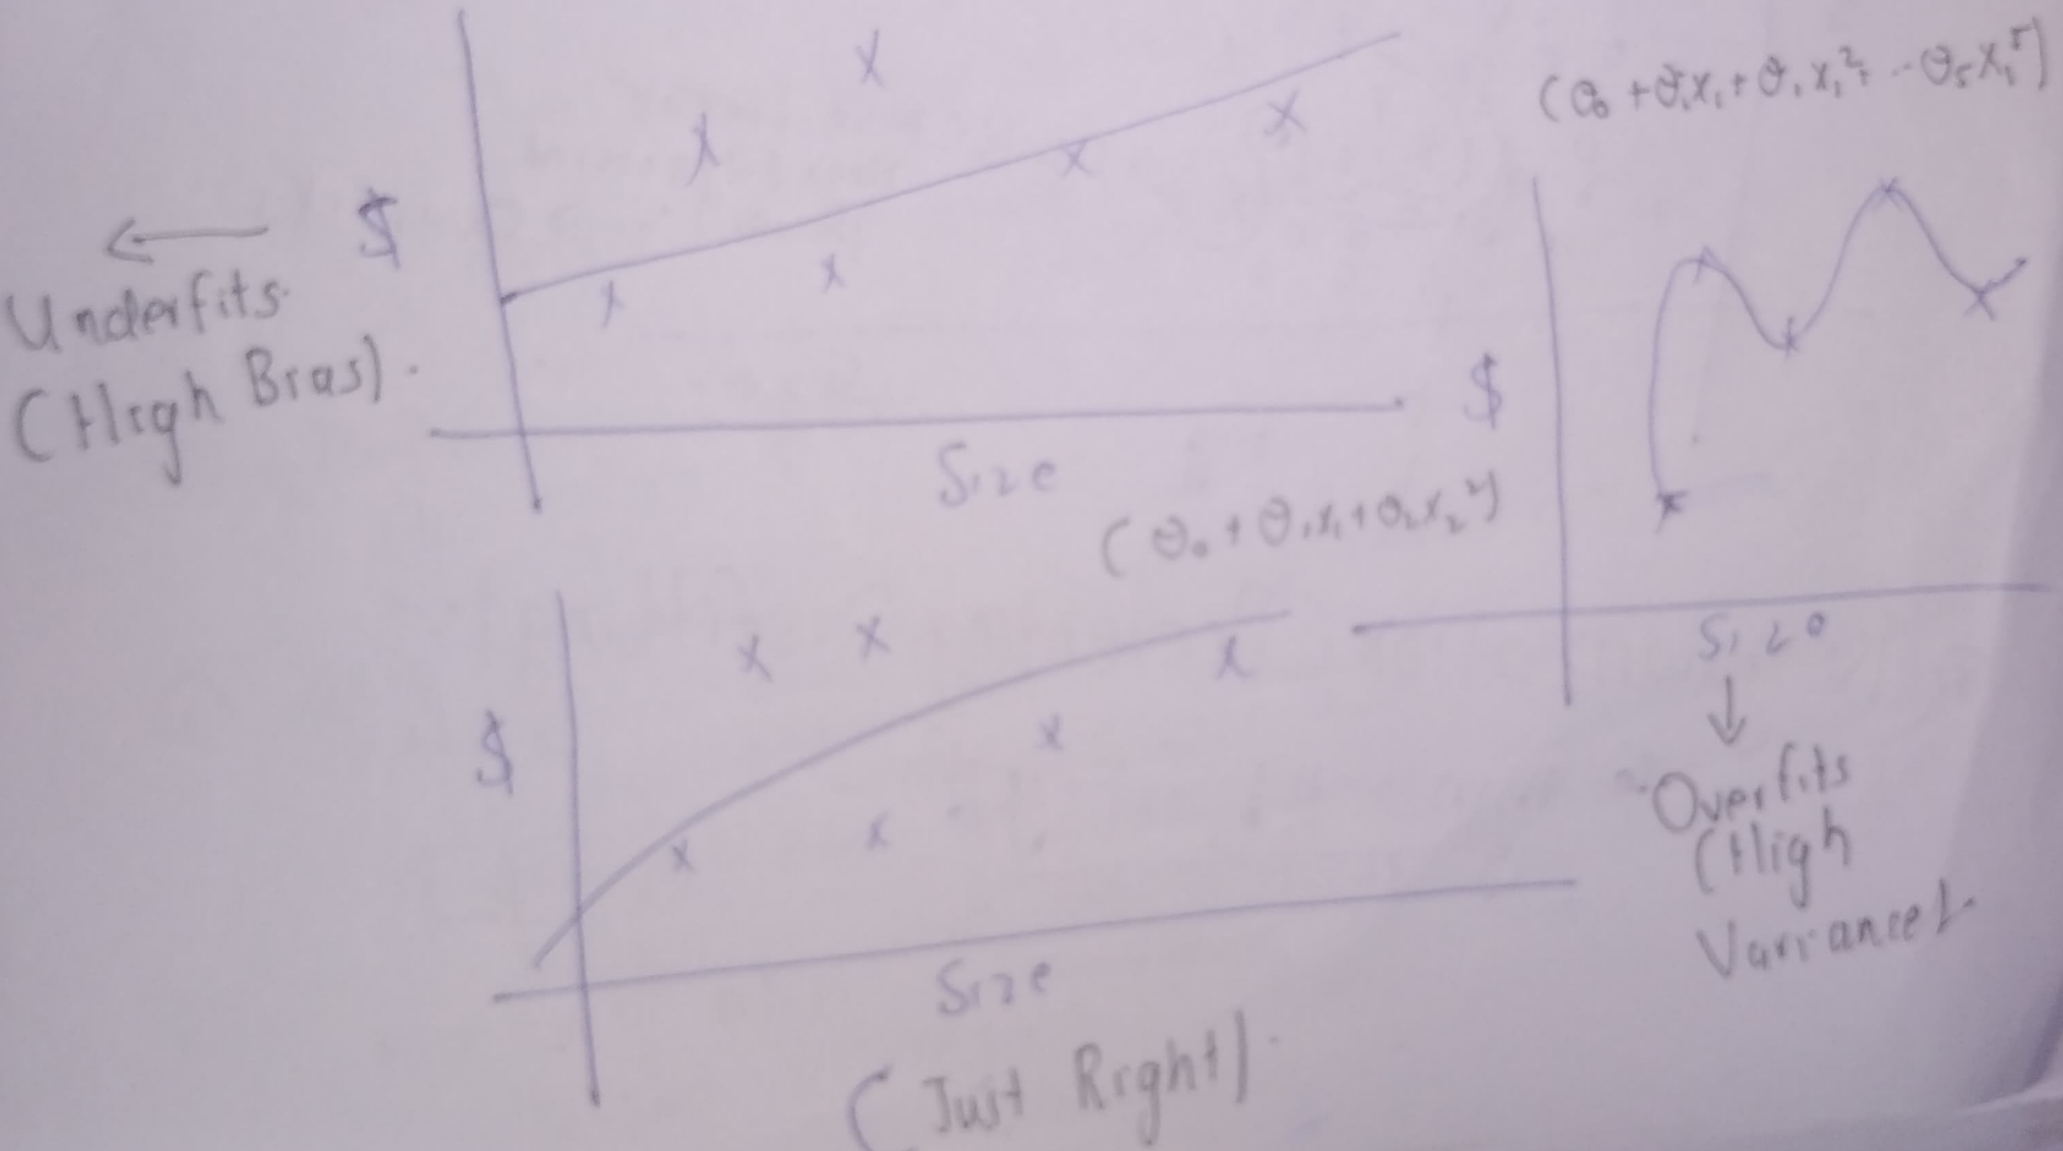

(*) The term high bias comes from the fact that our model was pr
Supposdely bias towards less complexity (linear) while actual data distribution was more complex (Quadratic). (High Bias).

(*) The term high ~~bias~~ Variance mean that everytime we fit our model on ~~stigth~~ slightly different data. our model would change to a large extent. Therefore our model fitting and predicting is highly varied given even a slight change in the distribution of data. (High Variance)

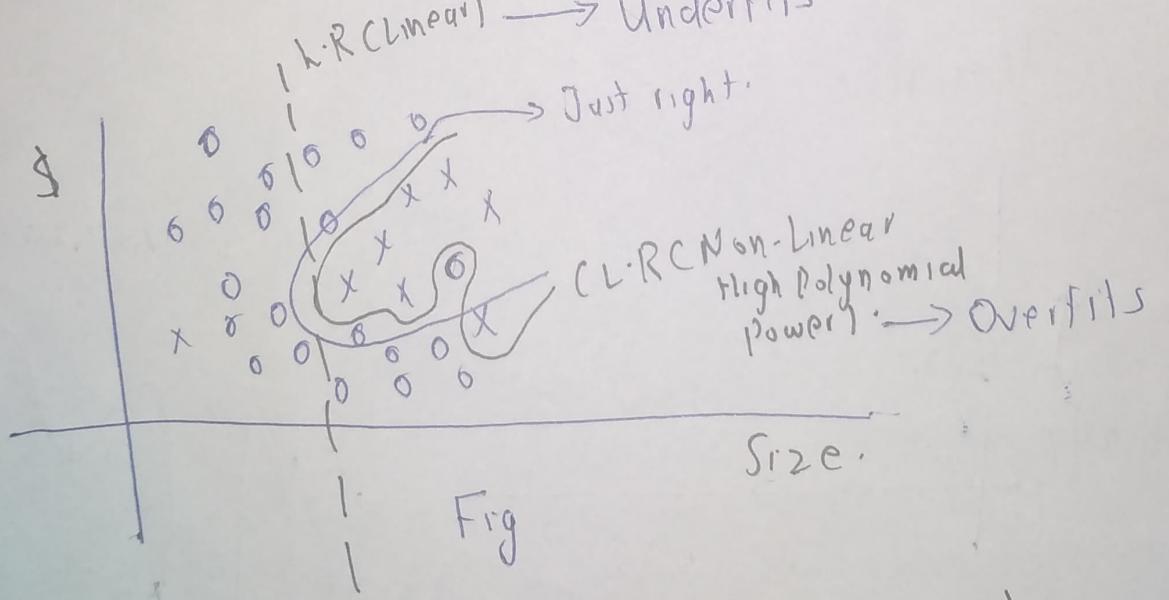

Fig

(*) Regularization (Preventing Overfitting):-

(a) ~~min~~ $\frac{1}{2} \| h_\vartheta(x) - y \|^2$
ϑ 2

$$\min_{\theta} \frac{1}{2} \sum_{i=1}^{m} \| y^{(i)} - \theta^T x^{(i)} \|^2 + \frac{\lambda}{2} \| \theta \|^2$$

Cost

Regularization term.

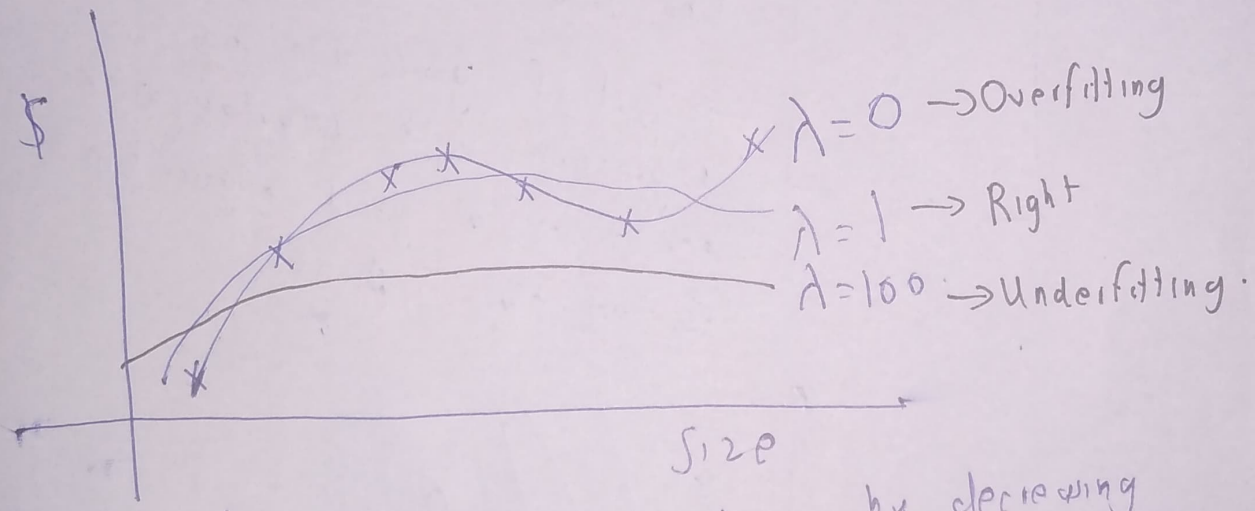

↳ Overfitting preventing by decreasing the size of θ.

↳ Very large λ mean θ will be forced to be close to zero.

~~(x) $\max_{\theta} \sum_{i=1}^{m}$ to~~

(x) $\arg\max_{\theta} \sum_{i=1}^{n} \log P(y^{(i)} | x^{(i)}; \theta) - \lambda \| \theta \|^2$.

↳ SVM does not overfit badly because $\min \|w\|^2$ objective function have same effective as regularization.

(*) Text classification

⊙ $m = 100$

$n = 10,000$

$x = \begin{bmatrix} 1 \\ 0 \\ \vdots \\ 1 \\ 0 \\ \vdots \end{bmatrix}$ a action

→ if you fit logistic regression to such a data where $n >> m$ then it will overfit the data ~~because~~

→ But if you use logistic regression with regularization then it is a good algorithm for text classification.

(*)

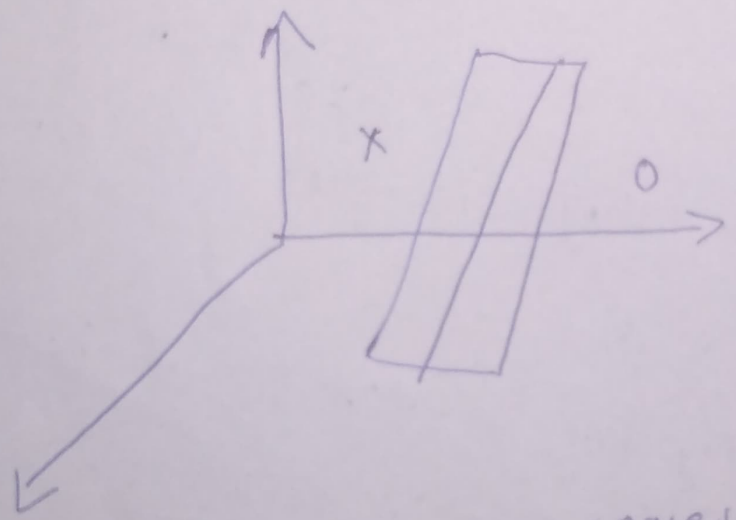

① More Example. mean logistic regression even without regularization won't overfit and would preform good provided $m > n$.

(*) Features at different scales should be normalized on the same range so that these learning models train faster and perform better.

→ Regularization and θ prior.

(*) $S = \{(x^{(i)}, y^{(i)})\}_{i=1}^{m}$

$$P(\theta|S) = \frac{P(S|\theta)\,P(\theta)}{P(S)}$$

(*) $\arg\max_{\theta} P(\theta|S) = \arg\max_{\theta} P(S|\theta)\cdot P(\theta)$

$$= \arg\max_{\theta} \underbrace{\left(\prod_{i=1}^{m} P(y^{(i)}|x^{(i)};\theta)\right)}_{\text{Generalized Linear model}} \cdot P(\theta).$$

$P(\theta): \theta \sim \mathcal{N}(0, \tau^2 I)$

$$P(\theta) = \frac{1}{\sqrt{2\pi}(\tau^2 I)^{1/2}} \exp\left(-\frac{1}{2}\theta^T (\tau^2 I)^{-1}\theta\right).$$

(*) If this is the prior distribution for θ and you plug it in then it is equivalent to regularization.

(*) Regularization is equivalent to assuming a prior θ Gaussian distribution.

Frequenist vs Bayesian Approach:

→ In Frequenist Approach we want to find the value of θ which makes the data as likely as possible

$\arg\max_{\theta} P(S|\theta)$ — MLE (Maximum Likelihood Estimate)

→ In Bayesian approach you already have some prior belief about θ in form of $P(\theta)$ (Gaussian prior)

↳ Our goal is to find value $\arg\max_{\theta} P(\theta|S)$ → MAP (Maximum a Posteriori)

↳ Adding Regularization

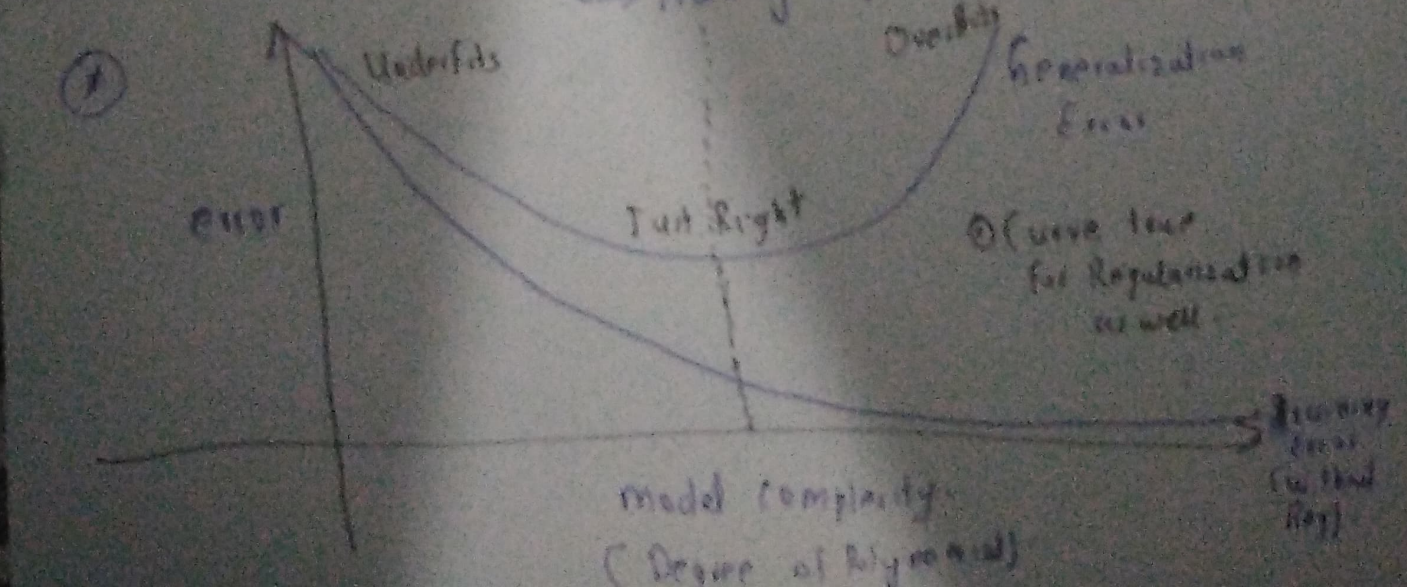

Training / Dev / Test

→ Spliting the Dataset into different subset:

→ 10000 examples

Different Hyperparameter needed to be choose ←

- Eg $\theta_0 + \theta_1 x$
- $\theta_0 + \theta_1 x + \theta_2 x^2$
- or choosing value λ
- or τ

→ Split your training data:

① S_{train} · S_{dev} (development) · S_{test}

② Train each model (option for different hyperparameters on S_{train}.
↳ Get some hypothesis h.

③ Measure error on S_{dev} for new developed hypothesis h.

④ Pick h with lowest error on S_{dev}.

⑤ If you want to check the unbiased performance of your model then evaluate on S_{test} since S_{dev} has already been optimized upon by your model.

(*) S_{test} gives the real performance of your model. Since S_{dev} is a biased estimate as model hyperparameter are adapted to perform on S_{dev}.

(*)

degree	S_{dev} Error	S_{test} Error
1	10	10
2	51	5
3	5·0	5·0
4	49 → Biased Estimate	5·0 → Actual Performance.
5	47	7
6	10	10
	⋮	⋮

(*) Sometimes you can ignore S_{test} in real world product but in publishing paper you need to have S_{test}.

(*) How much data to each split?

→ Historical:- (Percent thumb rule if dataset is small eg thousands examples)

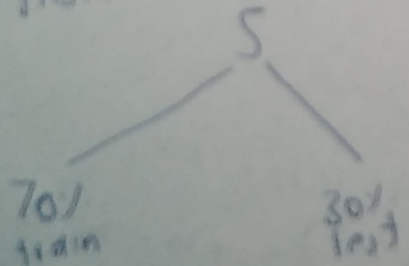

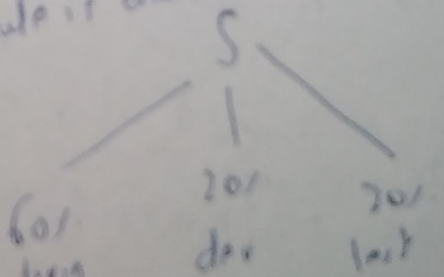

(*) Era of Big-data (10 million example):-

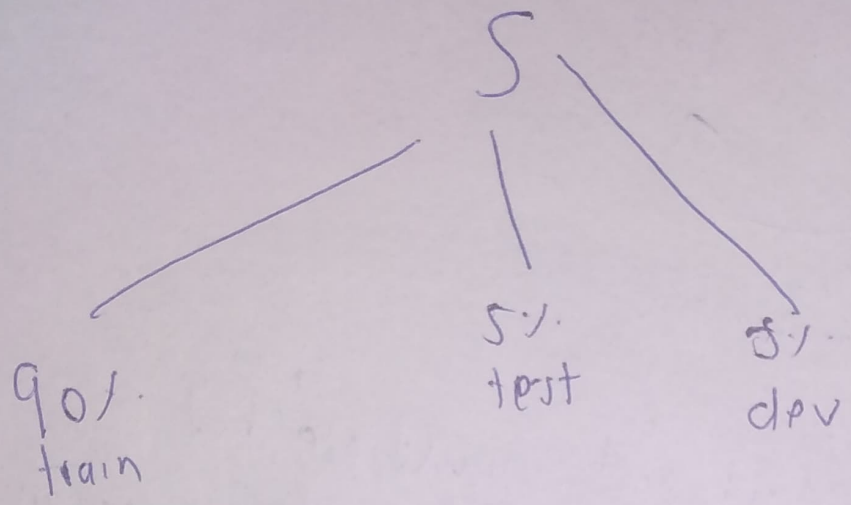

(*)

(*) But, Sometimes you need a largest test dataset if you the performance improvement between models is really small.

└ (*) Choose test and dev big enough to make meaningful comparision between different models.

(*) This Process of holding dev set from train set is called hold-out cross-validation (Simple hold-out cross validation).

└> Development set = Cross validation set (Dev)

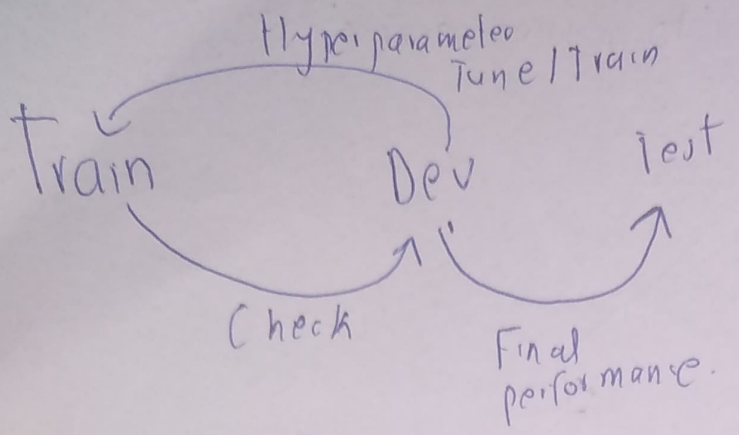

Fig.

→ Can evaluate on test set multiple times but don't make decision or tune hyper-parameter based on the performance of test set which is supposed to be an unbiased estimator benchmark
↳ Dev set evaluation is used for hyper-parameter tunning.

Small datasets:- (K-Fold-CV-(Cross-Validation).

- m=100 (Expensive Examples).
  70% — S_{train}, 30% — S_{dev}.
- This next procedure should be used only with small dataset:-
- K-Fold CV (Cross-Validation)
  Divide/Split data into five different subset

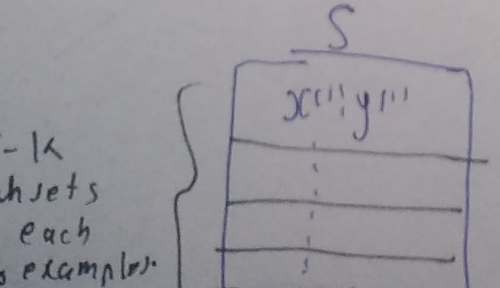

For i, ... k:-
Train on k-1 subsets.
Test on remaining 1 piece.

→ You will cross-validation for different hyperparameter and calculate average of metric for the cross-validations and then see which hyper-parameter performed the best.

→ Final Optional Step:-

→ Refit the Model with chosen hyperparameter on Complete dataset.

→ This procedure make efficent use of data but is computationally very expensive.

⊛ Extreme version of K-Fold CV:-
(For really small dataset):-

m=do (do example)

⊙ Leave-one-out CV:-

k=m → Divide your data into as many pieces as your training instances and leave only instance for evaluation at each iteration of validation.

(*) Feature Selection:-

- Small subset of features that are most useful for your task
- Most Important
- Depends on domain & context.

① Foward Search.

→ Start with $\mathcal{F} = \emptyset$

Repeat {

1) Try adding each feature i to $\mathcal{F}$, which single feature addition most improve dev set performance

2) Add that feature to $\mathcal{F}$ }

Example

- $X_1 \cdots X_5$ (Five Features)

$\emptyset$ $\quad h(x) = \theta_0$ (No feature)

$$\begin{bmatrix} \emptyset + x_1 \\ \emptyset + x_2 \longrightarrow h_\theta(x) = \theta + \theta_2 x_2 \quad \mathcal{F} = \{x_2\}. \\ \vdots \\ \emptyset + x_5 \end{bmatrix}$$

$$\begin{bmatrix} x_2 \; x_1 \longrightarrow \theta + \theta_1 x_1 + \theta_2 x_2 \quad \mathcal{F} = \{x_1, x_2\}. \\ x_2 \; x_3 \\ x_2 \; x_4 \\ x_2 \; x_5 \end{bmatrix}$$

→ The above example shows that you start with empty array of features and then at first iteration, train five different model with only one feature respectively. See which model with one specific feature perform the best and add that feature to your feature array

→ Then at second iteration you add another feature one at a time to previously choosen feature and see which model perform best with these two feature (one → new, one → previously choosen)

↳ You continue until you have either added all features or performance jump by new features addition is minimal.

⊛ Backward → Start with all features and remove features to one by one to see the performance degradation.

$$J = (X_1, X_2, X_3, X_4, X_5)$$

Repeat {

1) Remove Feature ith and train model with remaining features

2) Remove feature with most minimal decrease in performance when removed }.

Example

$Z = (X_1, \ldots X_5)$ (Five Features).

$[X_1, X_2, X_3, X_4]$
$[X_2, X_3, X_4, X_5]$
$[X_3, X_1, X_2, X_5]$ → Five different model and see whose accuracy is still the greatest and remove that feature (Say feature 5).
$[X_1\ X_2\ X_4\ X_5]$
$[X_2\ X_3\ X_1\ X_5]$

$Z = (X_1, X_2 X_3, X_4)$.

$[X_1, X_2, X_3]$
$[X_2\ X_3\ X_4]$
$[X_1\ X_2\ X_4]$ → Repeat until one feature is left or performance decrease is same across all features
$[X_2\ X_3\ X_4]$